

Lec 05 - NumPy Basics

Statistical Computing and Computation

Sta 663 | Spring 2022

Dr. Colin Rundel

NumPy Basics

What is NumPy?

NumPy is the fundamental package for scientific computing in Python. It is a Python library that provides a multidimensional array object, various derived objects (such as masked arrays and matrices), and an assortment of routines for fast operations on arrays, including mathematical, logical, shape manipulation, sorting, selecting, I/O, discrete Fourier transforms, basic linear algebra, basic statistical operations, random simulation and much more.

```
import numpy as np
```

Arrays

In general NumPy arrays are constructed from sequences (e.g. lists), nesting as necessary for the number of desired dimensions.

```
np.array([1,2,3])  
  
## array([1, 2, 3])  
  
np.array([[1,2],[3,4]])  
  
## array([[1, 2],  
##         [3, 4]])  
  
np.array([[[1,2],[3,4]], [[5,6],[7,8]]])  
  
## array([[[1, 2],  
##         [3, 4]],  
##         [[5, 6],  
##         [7, 8]]])
```

Note that NumPy stores data in row major order.

Some properties of arrays:

- Arrays have a fixed size at creation
- All data must be homogeneous (consistent type)
- Built to support vectorized operations
- Avoids copying whenever possible

```
np.array([True, False])
```

```
## array([ True, False])
```

```
np.array(["abc", "def"])
```

dtype

NumPy arrays will have a specific type used for storing their data, called their `dtype`. This is accessible via the `.dtype` attribute and can be set at creation using the `dtype` argument.

```
np.array([1,1]).dtype
```

```
## dtype('int64')
```

```
np.array([1.1, 2.2]).dtype
```

```
## dtype('float64')
```

```
np.array([True, False]).dtype
```

```
## dtype('bool')
```

```
np.array([1,2,3], dtype = np.uint8)
```

```
## array([1, 2, 3], dtype=uint8)
```

```
np.array([1,2,1000], dtype = np.uint8)
```

```
## array([ 1, 2, 232], dtype=uint8)
```

```
np.array([3.14159, 2.33333], dtype = np.double)
```

```
## array([3.14159, 2.33333])
```

```
np.array([3.14159, 2.33333], dtype = np.float16)
```

```
## array([3.14 , 2.334], dtype=float16)
```

See [here](#) for a list of dtypes and [here](#) for a more detailed description of how they are implemented.

Creating 1d arrays

Some common tools for creating useful 1d arrays:

```
np.arange(10)
```

```
## array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
np.arange(3, 5, 0.25)
```

```
## array([3. , 3.25, 3.5 , 3.75, 4. , 4.25, 4.5
```

```
np.linspace(0, 1, 11)
```

```
## array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7,
```

```
np.logspace(0, 2, 11)
```

```
## array([ 1.          ,  1.58489319,  2.51188643
##        , 6.30957344, 10.          , 15.84893192
##        , 39.81071706, 63.09573445, 100.
```

```
np.ones(4)
```

```
## array([1., 1., 1., 1.])
```

```
np.zeros(6)
```

```
## array([0., 0., 0., 0., 0., 0.])
```

```
np.full(3, False)
```

```
## array([False, False, False])
```

```
np.empty(4)
```

```
## array([1., 1., 1., 1.])
```

For the full list of creation functions see [here](#)

Creating 2d arrays (matrices)

Many of the same functions exist with some additional useful tools for common matrices,

```
np.eye(3)
```

```
## array([[1., 0., 0.],  
##         [0., 1., 0.],  
##         [0., 0., 1.]])
```

```
np.identity(2)
```

```
## array([[1., 0.],  
##         [0., 1.]])
```

```
np.zeros((2,2))
```

```
## array([[0., 0.],  
##         [0., 0.]])
```

```
np.diag([3,2,1])
```

```
## array([[3, 0, 0],  
##         [0, 2, 0],  
##         [0, 0, 1]])
```

```
np.tri(3)
```

```
## array([[1., 0., 0.],  
##         [1., 1., 0.],  
##         [1., 1., 1.]])
```

```
np.triu(np.full((3,3),3))
```

```
## array([[3, 3, 3],  
##         [0, 3, 3],  
##         [0, 0, 3]])
```

The NumPy documentation references a `matrix` class and related functions - this is no longer recommended, use the `ndarray` class instead.

Creating nd arrays

For higher dimensional arrays just add dimensions when constructing,

```
np.zeros((2,3,2))
```

```
## array([[[0., 0.],  
##         [0., 0.],  
##         [0., 0.]],  
##        [[0., 0.],  
##         [0., 0.],  
##         [0., 0.]])
```

```
np.ones((2,3,2,2))
```

```
## array([[[[1., 1.],  
##          [1., 1.]],  
##          [[1., 1.],  
##           [1., 1.]]],  
##         [[[1., 1.],  
##          [1., 1.]],  
##          [[1., 1.],  
##           [1., 1.]]],  
##         [[[1., 1.],  
##          [1., 1.]],  
##          [[1., 1.],  
##           [1., 1.]]]])
```

Subsetting

Arrays are subsetted using the standard python syntax with either indexes or slices, different dimensions are separated by commas.

```
x = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])  
x
```

```
## array([[1, 2, 3],  
##         [4, 5, 6],  
##         [7, 8, 9]])
```

```
x[0]
```

```
## array([1, 2, 3])
```

```
x[0,0]
```

```
## 1
```

```
x[0][0]
```

```
## 1
```

```
x[0:3:2, :]
```

Views and copies

Basic subsetting of ndarray objects does not result in a new object, but instead a "view" of the original object. There are a couple of ways that we can investigate this behavior,

```
x = np.arange(10)
y = x[2:5]
z = x[2:5].copy()
```

```
print("x =", x, ", x.base =", x.base)
## x = [0 1 2 3 4 5 6 7 8 9] , x.base = None

print("y =", y, ", y.base =", y.base)
## y = [2 3 4] , y.base = [0 1 2 3 4 5 6 7 8 9]

print("z =", z, ", z.base =", z.base)
## z = [2 3 4] , z.base = None

type(x); type(y); type(z)
## <class 'numpy.ndarray'>
```

```
np.shares_memory(x,y)
```

```
## True
```

```
np.shares_memory(x,z)
```

```
## False
```

```
np.shares_memory(y,z)
```

```
## False
```

```
y.flags
```

```
## C_CONTIGUOUS : True
```

Subsetting with ...

There is some special syntax available using ... which expands to the number of : needed to account for all dimensions,

```
x = np.arange(16).reshape(2,2,2,2)
x
```

```
## array([[[[ 0,  1],
##           [ 2,  3]],
##           [[ 4,  5],
##           [ 6,  7]]],
##           [[[ 8,  9],
##             [10, 11]],
##             [[12, 13],
##             [14, 15]]]])
```

```
x[0, 1, ...]
```

```
## array([[4, 5],
##         [6, 7]])
```

```
x[..., 1]
```

```
x[0, 1, :, :, :]
```

```
## array([[4, 5],
##         [6, 7]])
```

```
x[:, :, :, :, 1]
```

Subsetting with tuples

Unlike lists, an ndarray can be subset by a tuple containing integers,

```
x = np.arange(6)
x

## array([0, 1, 2, 3, 4, 5])

x[(0,1,3),]

## array([0, 1, 3])

x[(0,1,3)]

## Error in py_call_impl(callable, dots$args, dots
##
## Detailed traceback:
##   File "<string>", line 1, in <module>
```

More next time on why `x[(0,1,3)]` does not work.

```
x = np.arange(16).reshape(4,4)
x
```

```
## array([[ 0,  1,  2,  3],
##         [ 4,  5,  6,  7],
##         [ 8,  9, 10, 11],
##         [12, 13, 14, 15]])
```

```
x[(0,1,3), :]
```

```
## array([[ 0,  1,  2,  3],
##         [ 4,  5,  6,  7],
##         [12, 13, 14, 15]])
```

```
x[:, (0,1,3)]
```

```
## array([[ 0,  1,  3],
##         [ 4,  5,  7],
##         [ 8,  9, 11],
##         [12, 13, 15]])
```

Subsetting assignment

Most of the subsetting approaches we've just seen can also be used for assignment, just keep in mind that we cannot change the size or type of the ndarray,

```
x = np.arange(9).reshape((3,3)); x
```

```
## array([[0, 1, 2],  
##         [3, 4, 5],  
##         [6, 7, 8]])
```

```
x[0,0] = -1  
x
```

```
## array([[-1,  1,  2],  
##         [ 3,  4,  5],  
##         [ 6,  7,  8]])
```

```
x[0, :] = -2  
x
```

```
## array([[-2, -2, -2],  
##         [ 3,  4,  5],  
##         [ 6,  7,  8]])
```

```
x[0:2,1:3] = -3  
x
```

```
## array([[-2, -3, -3],  
##         [ 3, -3, -3],  
##         [ 6,  7,  8]])
```

```
x[(0,1,2), (0,1,2)] = -4  
x
```

```
## array([[-4, -3, -3],  
##         [ 3, -4, -3],  
##         [ 6,  7, -4]])
```

Reshaping arrays

The dimensions of an array can be retrieved via the `.shape` attribute, these values can be changed by

```
x = np.arange(6)  
x
```

```
## array([0, 1, 2, 3, 4, 5])
```

```
y = x.reshape((2,3))  
y
```

```
## array([[0, 1, 2],  
##         [3, 4, 5]])
```

```
np.shares_memory(x,y)
```

```
## True
```

```
z = x  
z.shape = (2,3)  
z
```

```
## array([[0, 1, 2],  
##         [3, 4, 5]])
```

```
x
```

```
## array([[0, 1, 2],  
##         [3, 4, 5]])
```

```
np.shares_memory(x,z)
```

Implicit dimensions

When reshaping an array, the value -1 can be used to automatically calculate a dimension,

```
x = np.arange(6)  
x  
## array([0, 1, 2, 3, 4, 5])  
x.reshape((2,-1))
```

```
## array([[0, 1, 2],  
##         [3, 4, 5]])
```

```
x.reshape((-1,3,2))
```

```
## array([[[0, 1],  
##          [2, 3],  
##          [4, 5]]])
```

```
x.reshape(-1)
```

```
## array([0, 1, 2, 3, 4, 5])
```

```
x.reshape((-1,4))
```

Flattening arrays

We've just seen the most common approach to flattening an array (`.reshape(-1)`), there are two additional methods / functions:

- `ravel` which creates flattened view of the array and
- `flatten` which creates a flattened copy of the array.

```
x = np.arange(6).reshape((2,3))
x
```

```
## array([[0, 1, 2],
##         [3, 4, 5]])
```

```
y = x.ravel()
y
```

```
## array([0, 1, 2, 3, 4, 5])
np.shares_memory(x,y)
## True
```

```
z = x.flatten()
z
```

```
## array([0, 1, 2, 3, 4, 5])
np.shares_memory(x,z)
## False
```

Resizing

The size of an array cannot be changed but a new array with a different size can be created from an existing array via the `resize` function and method. Note these have different behaviors around what values the new entries will have.

```
x = np.resize(np.ones((2,2)), (3,3))
x
```

```
## array([[1., 1., 1.],
##         [1., 1., 1.],
##         [1., 1., 1.]])
```

```
y = np.ones((2,2))
y.resize((3,3), refcheck=False)
y
```

```
## array([[1., 1., 1.],
##         [1., 0., 0.],
##         [0., 0., 0.]])
```

The `refcheck` argument should not generally be needed, here it is a side effect of using `reticulate` (I think).

Joining arrays

`concatenate()` is a general purpose function for this, with specialized versions `hstack()`, `vstack()`, and `dstack()` for rows, columns, and slices respectively.

```
x = np.arange(4).reshape((2,2)); x
```

```
## array([[0, 1],  
##         [2, 3]])
```

```
np.concatenate((x,y), axis=0)
```

```
## array([[0, 1],  
##         [2, 3],  
##         [4, 5],  
##         [6, 7]])
```

```
np.concatenate((x,y), axis=1)
```

```
## array([[0, 1, 4, 5],  
##         [2, 3, 6, 7]])
```

```
np.concatenate((x,y), axis=2)
```

```
## Error in py_call_impl(callable, dots$args, dots$keywords): Ax  
##  
## Detailed traceback:  
See block(s) for a more efficient approach to constructing block matrices.  
## File "<string>" line 1 in <module>  
## File "<__array_function__ internals>", line 180, in concat
```

```
y = np.arange(4,8).reshape((2,2)); y
```

```
## array([[4, 5],  
##         [6, 7]])
```

```
np.vstack((x,y))
```

```
## array([[0, 1],  
##         [2, 3],  
##         [4, 5],  
##         [6, 7]])
```

```
np.hstack((x,y))
```

```
## array([[0, 1, 4, 5],  
##         [2, 3, 6, 7]])
```

```
np.dstack((x,y))
```

```
## array([[[0, 4],  
##                 [1, 5]],  
##                 [[2, 6],  
##                 [3, 7]]])
```

NumPy numerics

Basic operators

All of the basic mathematical operators in Python are implemented for arrays, they are applied element-wise to the array values.

```
np.arange(3) + np.arange(3)
```

```
## array([0, 2, 4])
```

```
np.arange(3) - np.arange(3)
```

```
## array([0, 0, 0])
```

```
np.arange(3) + 2
```

```
## array([2, 3, 4])
```

```
np.full((2,2), 2) ** np.arange(4).reshape((2,2))
```

```
## array([[1, 2],  
##        [4, 8]])
```

```
np.full((2,2), 2) ** np.arange(4)
```

```
np.arange(3) * np.arange(3)
```

```
## array([0, 1, 4])
```

```
np.arange(1,4) / np.arange(1,4)
```

```
## array([1., 1., 1.])
```

```
np.arange(3) * 3
```

```
## array([0, 3, 6])
```

Mathematical functions

The package provides a wide variety of basic mathematical functions that are vectorized, in general they will be faster than their base equivalents (e.g. `np.sum()` vs `sum()`),

```
np.sum(np.arange(1000))  
## 499500  
  
np.cumsum(np.arange(10))  
## array([ 0,  1,  3,  6, 10, 15, 21, 28, 36, 45])  
  
np.log10(np.arange(1,11))  
## array([0.        , 0.30103    , 0.47712125, 0.60205999, 0.69897    ,  
##        0.77815125, 0.84509804, 0.90308999, 0.95424251, 1.        ])  
  
np.median(np.arange(10))  
## 4.5
```

Matrix multiplication

Is supported using the `matmul()` function or the `@` operator,

```
x = np.arange(6).reshape(3,2)
y = np.tri(2,2)
```

```
x @ y
```

```
## array([[1., 1.],
##         [5., 3.],
##         [9., 5.]])
```

```
y.T @ y
```

```
## array([[2., 1.],
##         [1., 1.]])
```

```
np.matmul(x.T, x)
```

```
## array([[20, 26],
##         [26, 35]])
```

```
y @ x
```

Other linear algebra functions

All of the other common linear algebra functions are (mostly) implemented in the `linalg` submodule. See [here](#) for more details.

```
np.linalg.det(y)
## 1.0

np.linalg.eig(x.T @ x)
## (array([ 0.43988174, 54.56011826]), array([[[-0.79911221, -0.6011819 ],
## [ 0.6011819 , -0.79911221]]))

np.linalg.inv(x.T @ x)
## array([[ 1.45833333, -1.08333333],
##        [-1.08333333,  0.83333333]])

np.linalg.cholesky(x.T @ x)
## array([[4.47213595, 0.          ],
##        [5.81377674, 1.09544512]])
```

Random values

NumPy has another submodule called `random` for functions used to generate random values, In order to use this, you should construct a generator via `default_rng()`, with or without a seed, and then use the generator's methods to obtain your desired random values.

```
rng = np.random.default_rng(seed = 1234)

rng.random(3) # ~ Uniform [0,1]

## array([0.97669977, 0.38019574, 0.92324623])

rng.normal(0, 2, size = (2,2))

## array([[ 0.30523839,  1.72748778],
##        [ 5.82619845, -2.95764672]])

rng.binomial(n=5, p=0.5, size = 10)

## array([2, 4, 2, 2, 3, 4, 4, 3, 3, 3])
```

Example - Linear regression with NumPy